

CSE8803/CX4803

Machine Learning in Computational Biology

Lecture 7:
PCA, Autoencoder, VAE

Yunan Luo

Logistics: presentation date

- Presentation team finalized (see Canvas -> Quizzes)
- Each team to submit **3 preferred dates** via a Google form
 - Link on Ed
- Deadline: **Feb 7 (Mon), 11:59 PM**
- Distribute **10 points** into the 3 dates
- Detailed announcement on Ed
- Paper selection
 - A paper list will be released
 - More details to follow

Bidding for presentation dates

You will select from the following dates:

2/21, 2/23, 3/7, 3/16, 4/4, 4/6, 4/11, 4/13, 4/18, 4/20, 4/25

You have 10 points in total. Please distribute them to at most 3 dates of your choice. Please write the dates in the same format as shown above (dd/mm). The number of points for each date does not have to be an integer.

Deadline is Monday 2/7, 11:59pm.

 **luoyunan@gmail.com** (not shared) [Switch account](#)



* Required

Your name *

Your answer

Your Group ID *

Your answer

Your 1st choice date *

Your answer

Number of points for your 1st choice date *

Your answer

Presenting the paper

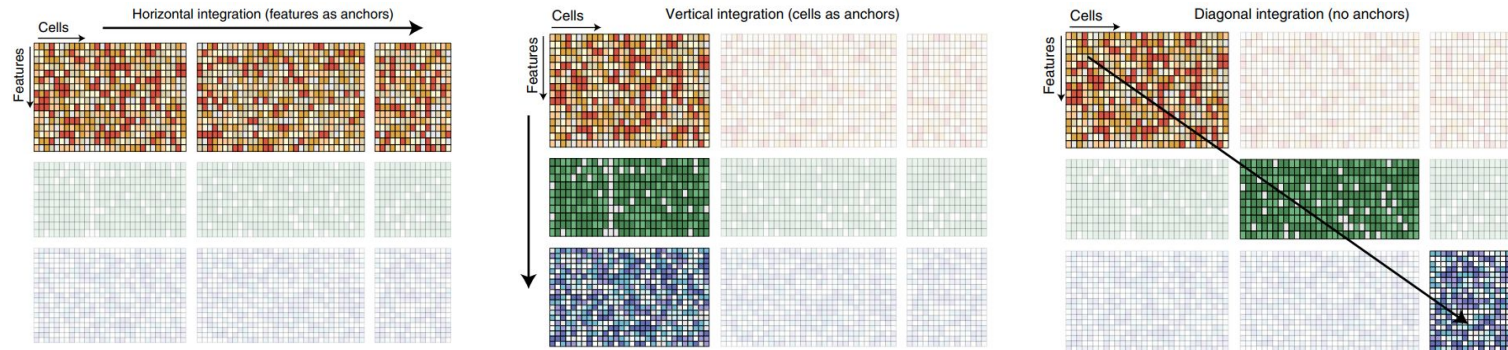
- Briefly introduce the biological problem
 - E.g., why it's important? What's the motivation? Why use ML/computation methods?
- Clearly state the computational problem
 - E.g., what's are the input and output? What does the data look like?
- Present the methods and results
 - Try to identify key ideas in methods and key takeaways from results
- Optional
 - Your comments on the presented paper
 - Survey of related work
 - Follow-up research ideas or applications

Optional: share your comments on the paper

- Consider you are a reviewer of the paper, share your critical (not necessarily negative) comments/review of the paper
 - Strengths and/or weakness
 - Novelty/significance of the contribution
 - Soundness of the evaluation
 - Further improvement of the paper

Optional: survey of related work

- Positioning the paper in the context of previous and subsequent work
 - Any prior papers that substantially influenced the presented paper?
 - Any newer papers that are largely built on the presented paper
 - Make a brief comparison of them in terms of motivation, strength, limitation, methodology, etc
- Example: single-cell data integration methods



Optional: possible follow-up projects or applications

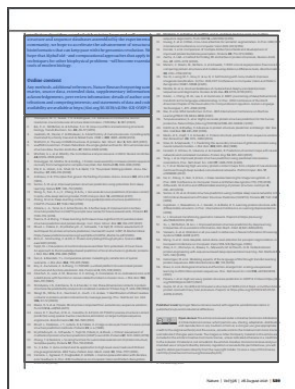
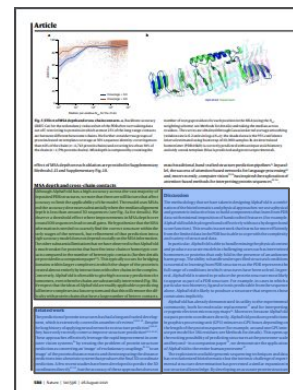
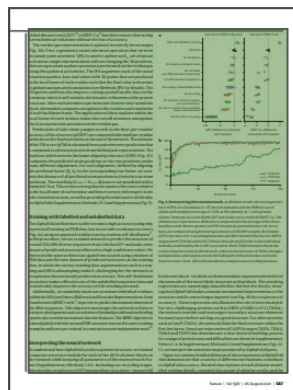
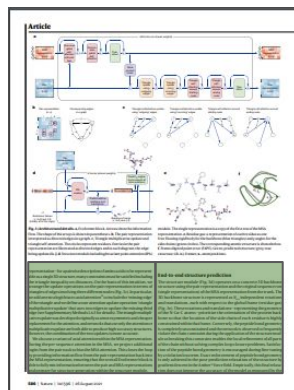
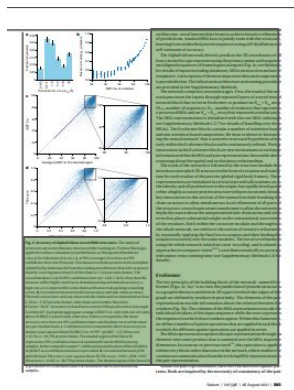
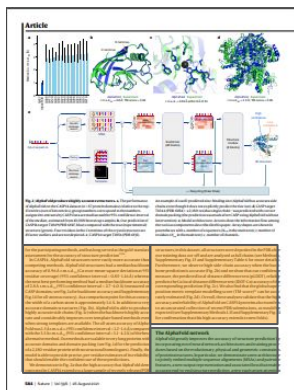
- Propose a follow-up research project idea
 - Any improvements of the proposed method in the paper?
 - Can you think about a solution that addresses some limitations of the proposed method?
- Propose an application based on the paper
 - Think about a new application in biology (not discussed in the paper or our class yet) where the method can be applied to.

Suggestions

- Make sure the presentation ends within the time limit
- Don't put too much information on a single slide
 - Avoid using long sentences or dense tables
 - Use clear, short text
 - Use illustrations/demos to show the methods/data/results
- Connect the presented paper with what we have seen in the class, if possible
- Don't forget important information in the Supplementary Information of journal papers

Results

(may have a method overview)

[illegible]

Methods

Jumper, J., Evans, R., Pritzel, A. *et al.* Highly accurate protein structure prediction with AlphaFold. *Nature* 596, 583–589 (2021).

Supplementary information

<https://doi.org/10.1038/s41586-021-03819-2>

Highly accurate protein structure prediction with AlphaFold

In the format provided by the authors and unedited

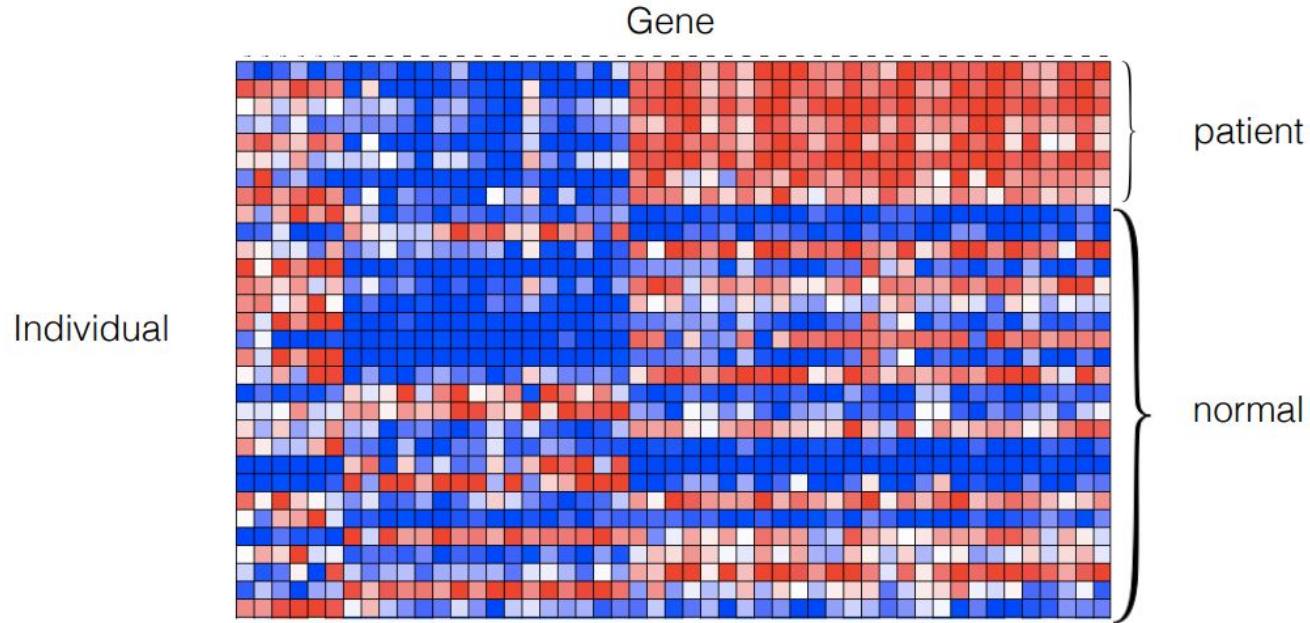
Full algorithm details
Supplementary Information (62 page)

Outline

- PCA
- Autoencoder
- Variational Autoencoder (VAE)

Week	Date	Topic	Contents	Instructor
1	1/10/2022	Introduction	Course intro & how to present papers	Zhang
	1/12/2022	Learning from sequence data	Dynamic programming & sequence alignment I	Zhang
2	1/17/2022		No class (MLK Day)	
	1/19/2022		Sequence alignment II	Zhang
3	1/24/2022		HMM & gene/motif finding	Zhang
	1/26/2022		HMM & Profile HMM	Zhang
4	1/31/2022		Deep learning for DNA/protein sequence	Luo
	2/2/2022	Learning from high-dim data	Learn from high-dim data: PCA, autoencoder & VAE	Luo
5	2/7/2022		Learn from high-dim data: MDS, tSNE, UMAP	Zhang
	2/9/2022		Clustering I	Zhang
6	2/14/2022		Clustering II	Zhang
	2/16/2022		Clustering III	Zhang
7	2/21/2022		Student presentation 1-3	
	2/23/2022		Student presentation 4-6	
8	2/28/2022	Learning from structure data	RNA structure prediction	Luo
	3/2/2022		Deep learning for structures (protein structure prediction)	Luo
9	3/7/2022	Learning from network data	Student presentation 7-9	
	3/9/2022		Network basics & traditional ML for graphs	Luo
10	3/14/2022		Network embeddings	Luo
	3/16/2022		Student presentation 10-12	
11	3/21/2022		No class (Spring Break)	
	3/23/2022		No class (Spring Break)	
12	3/28/2022		Graphical Models	Luo
	3/30/2022		Deep learning for networks (graph neural networks)	Luo

Gene expression matrix




$\text{dim}(\text{features}) \gg \text{num}(\text{samples})$

High-dimensional data


High-dimensional data

- Each sample has a large number of features/attributes



Why is high-dimension a problem? The curse of dimensionality:

- Volume of space increases exponentially so data becomes very sparse;  sparsity
- Increases the effort of searching drastically
- Makes it harder to calculate (accurate) distances between samples
- Redundancy of data
- A large number of training data samples is required to train a model for high-dim data
- Overfitting

Overfitting



$$p(\text{number of parameters}) \gg n(\text{number of data points})$$



A solution: dimensionality reduction

Benefits:

- Reduce redundancy of data
- Identify the most relevant information (find and filter noise) & Cleaning the data
- Reduce computational complexity & Speeding up subsequent learning task
- Building simpler model later
- Visualizing, exploring and understanding the data

Dimensionality reduction: approaches

- Linear transformation:
 - PCA
 - NMF
- Non-linear transformation
 - Autoencoder, VAE
 - MDS
 - tSNE
 - UMAP
- Different methods have different *objectives*



Principal Component Analysis (PCA)

Component analysis

How to understand the main signals from the data?

Key assumptions

1. Low-rank assumption: High-dimensional data lies on a lower dimensional space (a.k.a, manifold)
2. Projections in the lower-dimensional space describes major properties of the data



Principal Component Analysis (PCA)

Goal: Find a projection of the data onto directions that **maximize variance** of the original data

- **Intuition:** those are directions in which most information is encoded

Definition: Principal Components (PC) are orthogonal directions that capture most of the variance in the data




Figure credit: Le Song

PCA: Finding principal components

- 1st PC:
 - Projection of data points along 1st PC discriminates data most along any one direction
- 2nd PC:
 - Next orthogonal direction of greatest variability
- 3rd PC...

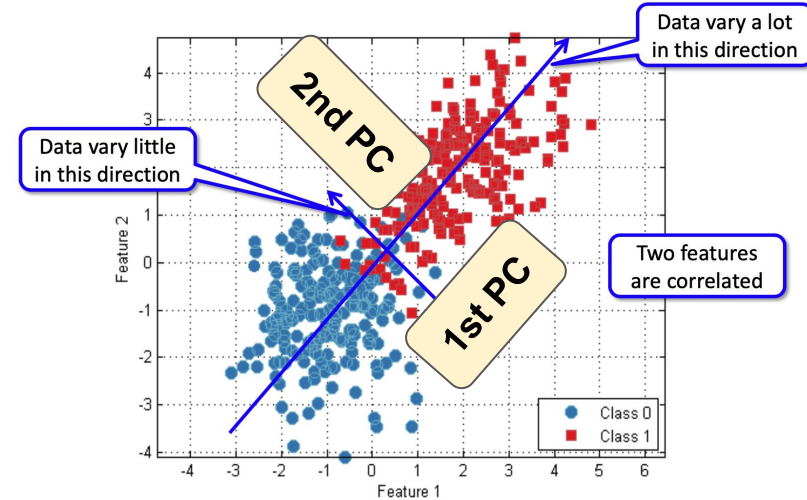


Figure credit: Le Song

PCA notation

- Input data points: matrix $X = [x_1, x_2, \dots, x_N]$ of size $D \times N$
- x_i is the i -th column, i.e., the i -th example
- x_{ij} is the j -th feature of example i
- We assume the data is **centered**, i.e., $\frac{1}{N} \sum_{i=1}^N x_i = \vec{0}$
 - If not centered, replace x_i by $x_i - \mu$, where $\mu = \frac{1}{N} \sum_{i=1}^N x_i$

Finding the 1st PC

Given N data points, $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N]_{D \times N}$, $\mathbf{x}_i \in \mathbb{R}^D$, find a direction \mathbf{w} where $\|\mathbf{w}\| = 1$, such that the variation of the data along direction \mathbf{w} is maximized.

- The sample variance on the projected on vector \mathbf{w} is $\sum_{i=1}^N (\mathbf{w}^T \mathbf{x}_i)^2 = \mathbf{w}^T \mathbf{X} \mathbf{X}^T \mathbf{w}$

Find the 1st PC by solving the following optimization problem

$$\begin{aligned} \max_{\mathbf{w}} \quad & \mathbf{w}^T \mathbf{X} \mathbf{X}^T \mathbf{w} \\ \text{such that:} \quad & \|\mathbf{w}\| = 1 \end{aligned}$$

Finding the 1st PC

$$\begin{aligned} \max_{\mathbf{w}} \quad & \mathbf{w}^T \mathbf{X} \mathbf{X}^T \mathbf{w} \\ \text{such that:} \quad & \|\mathbf{w}\| = 1 \end{aligned}$$

- Construct Lagrange multiplier

$$\max_{\mathbf{w}} \mathbf{w}^T \mathbf{X} \mathbf{X}^T \mathbf{w} - \lambda(\|\mathbf{w}\| - 1)$$

- Take the derivative with respect to \mathbf{w} and set it to 0 \Rightarrow solutions are vectors \mathbf{w} such that

$$\mathbf{X} \mathbf{X}^T \mathbf{w} = \lambda \mathbf{w}$$

----- Eigenvector of $\mathbf{X} \mathbf{X}^T$!

$$\mathbf{X} \mathbf{X}^T \mathbf{w} = \lambda \mathbf{w}$$

The eigenvalue problem

- For a given matrix \mathbf{A}

$$\mathbf{A} \mathbf{w} = \lambda \mathbf{w}$$

\mathbf{w} is the eigenvector and λ is the eigenvalue

- There are multiple solutions $\mathbf{w}_1, \mathbf{w}_2, \dots$,
with different (or same) eigenvalues $\lambda_1, \lambda_2, \dots$
- The eigenvectors are ortho-normal (symmetric,
positive semi-definite)
 - $\mathbf{w}_i^T \mathbf{w}_j = 0, \mathbf{w}_i^T \mathbf{w}_i = 1$
- Let $\mathbf{A} = \mathbf{X} \mathbf{X}^T$ and find the eigenvectors and eigenvalues of \mathbf{A}

PCA formally

- If we rank eigenvalues from large to small
 - The 1st PC is the eigenvector of \mathbf{XX}^T associated with the largest eigenvalue
 - The 2nd PC is the eigenvector of \mathbf{XX}^T associated with the 2nd largest eigenvalue
 - ...
- The eigenvalue $\lambda_i / \sum \lambda_i$ denotes the percentage of variance accounted for by the i-th PC \mathbf{w}_i

Q1: how to find eigenvalues/eigenvectors?

Singular value decomposition (SVD)

The SVD is a factorization of a $m \times n$ matrix into

$$A = U \Sigma V^T$$

where U is a $m \times m$ orthogonal matrix, V^T is a $n \times n$ orthogonal matrix and Σ is a $m \times n$ diagonal matrix.

$$A = \begin{pmatrix} \vdots & \dots & \vdots \\ \mathbf{u}_1 & \dots & \mathbf{u}_n \\ \vdots & \dots & \vdots \end{pmatrix} \begin{pmatrix} \sigma_1 & & \\ & \ddots & \\ & & \sigma_n \end{pmatrix} \begin{pmatrix} \dots & \mathbf{v}_1^T & \dots \\ \vdots & \vdots & \vdots \\ \dots & \mathbf{v}_n^T & \dots \end{pmatrix}$$

SVD in Python:

```
from scipy import linalg  
U, s, Vh = linalg.svd(A)
```

Singular value decomposition (SVD)

Theorem: if a **square** matrix S is a **real** and **symmetric** matrix, then its SVD can be represented as


$$\mathbf{S} = \mathbf{V}\mathbf{\Lambda}\mathbf{V}^T$$

where columns of \mathbf{V} are **eigenvectors** of \mathbf{S} and diagonal elements of $\mathbf{\Lambda}$ are **eigenvalues** of \mathbf{S}

$$\mathbf{\Lambda} = \text{diag}(\lambda_1, \dots, \lambda_m), \quad \lambda_i \geq \lambda_{i+1}$$


Q2: how many PCs

- The eigenvalue λ_i denotes the amount of variability captured along dimension \mathbf{w}_i
- Can ignore the components of lower variance (less significant)



Alternative interpretation 1: residual minimization

PCA finds vectors \mathbf{v} such that projection on to these vectors minimizes reconstruction error



([image source](#))


Alternative interpretation 2: low-rank approximation

PCA seeks the best rank- k approximation to the matrix A in the least-squares sense, by solving

$$\begin{array}{ll}\text{minimize} & \|A - Z\|_F^2 \\ \text{subject to} & \text{Rank}(Z) \leq k,\end{array}$$

with variable $Z \in \mathbf{R}^{m \times n}$. Here, $\|\cdot\|_F$ is the Frobenius norm of a matrix, *i.e.*, the square root of the sum of the squares of the entries.

Example: Tissue-specific gene expression




Summary: PCA

- What you should know:
 - Goal: Find a **projection** of the data onto directions that **maximize variance** of the original data
 - Optimization objective & algorithm
- Pros
 - Eigenvector method
 - No tuning of parameters
 - No local optima
- Cons
 - Only based on covariance (2nd order statistics)
 - Limited to linear projections
- Next: Nonlinear dimensionality reduction

Autoencoder


Autoencoders

- A neural network to find latent space representation of the original data
 - Unsupervised method (with no labeled training data)
- To make this non-trivial, we add a **bottleneck layer** whose dimension is much smaller than the input




Why autoencoders?

- Map high-dimensional data to 2D for **visualization**
- Compression
- Learn abstract **features** in an unsupervised way so you can apply them to a supervised task
- Learn a semantically meaningful representation where you can, e.g., **interpolate** between different images.



Autoencoders: approach

- **Goal:** Find the **latent space representation** that *best* represent the important information in the original data
 - Recall PCA: maximize the variance
- **Approach:** bottleneck layer
 - Forces the network to create a compressed representation of the input data (dimensionality reduction)
 - Forces the network to remove redundancy and noise
- **Objective:** reconstruction error $\mathcal{L}(\mathbf{x}, \tilde{\mathbf{x}}) = \|\mathbf{x} - \tilde{\mathbf{x}}\|^2$
 - Can add regularization term to avoid overfitting (identity mapping)



Autoencoders: connection to PCA

Loss function: $\mathcal{L}(x, \tilde{x})$ (reconstruction error)

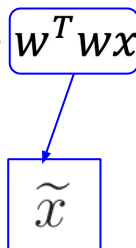
Mean square error (MSE):


$$\frac{1}{D} \sum_i \|x_i - \tilde{x}_i\|^2$$

What if we remove non-linearity in NN?

When we remove the non-linearity term in neural network (activation function), (and force encoder and decoder to have the same weights) autoencoder is equivalent to PCA :

$$\hat{w} = \arg \min_w \mathbb{E}[\|x - \boxed{w^T w x}\|^2]$$





PCA

$$\mathbf{Z}_{(rxN)} = \mathbf{W}_{(rxD)}^T \mathbf{X}_{(D \times N)}$$

A different objective function:

$$\begin{aligned} \min & \|\mathbf{X} - \mathbf{WZ}\|^2 \\ &= \min \|\mathbf{X} - \mathbf{W}\mathbf{W}^T\mathbf{X}\|^2 \end{aligned}$$

Autoencoders: connection to PCA

Autoencoders learn to project the data, not onto a subspace, but onto a nonlinear **manifold**

Linear vs nonlinear dimensionality reduction




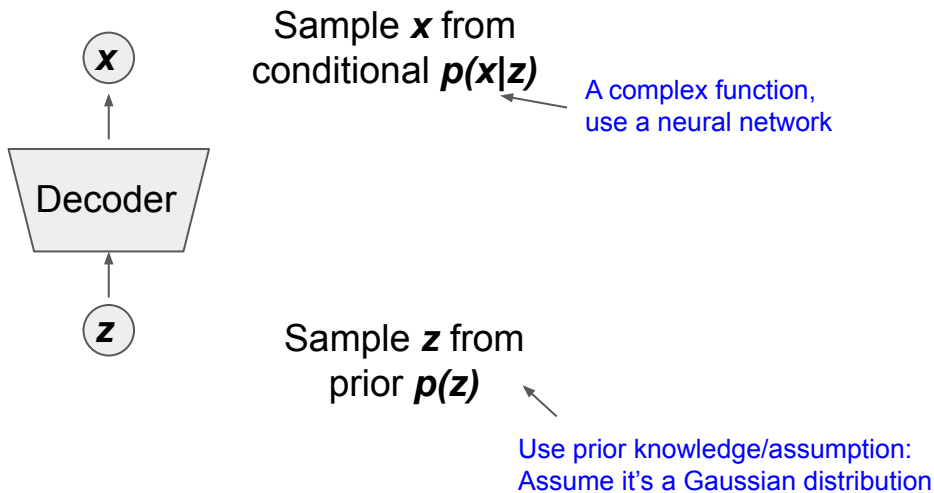
Image source: <https://www.jeremyjordan.me/autoencoders/>

Variational autoencoder (VAE)

Variational Autoencoders

Generative models: allow us to sample from the models to generate new data.

Assume training data \mathbf{x} is generated from underlying unobserved (latent) representation \mathbf{z}




Intuition: consider \mathbf{x} as an image, \mathbf{z} as latent factors to generate \mathbf{x} : attributes, orientation, etc

Generative models

- Maximum likelihood

$$p(\mathbf{x}) = \int p(\mathbf{z})p(\mathbf{x} | \mathbf{z}) d\mathbf{z}$$

- Problem of autoencoder
 - If \mathbf{z} is low-dimensional and the decoder is deterministic, then $p(\mathbf{x}) = 0$ almost everywhere
- Idea of VAE: instead of encoding an input as a single point, we encode it as a **distribution** over the latent space.



Maximum likelihood

$$p(\mathbf{x}) = \int p(\mathbf{z}) p(\mathbf{x} | \mathbf{z}) d\mathbf{z}$$

Intractable to
compute for
every \mathbf{z}



Gaussian
prior



Neural
network



Maximizing a lower bound

$$\begin{aligned}\log p(\mathbf{x}) &= \log \int p(\mathbf{z}) p(\mathbf{x}|\mathbf{z}) d\mathbf{z} \\ &= \log \int q(\mathbf{z}) \frac{p(\mathbf{z})}{q(\mathbf{z})} p(\mathbf{x}|\mathbf{z}) d\mathbf{z}\end{aligned}$$

Auxiliary distribution $q(\mathbf{z})$


Maximizing a lower bound

$$\begin{aligned}\log p(\mathbf{x}) &= \log \int p(\mathbf{z}) p(\mathbf{x}|\mathbf{z}) d\mathbf{z} \\ &= \log \int q(\mathbf{z}) \frac{p(\mathbf{z})}{q(\mathbf{z})} p(\mathbf{x}|\mathbf{z}) d\mathbf{z} \\ &\geq \int q(\mathbf{z}) \log \left[\frac{p(\mathbf{z})}{q(\mathbf{z})} p(\mathbf{x}|\mathbf{z}) \right] d\mathbf{z}\end{aligned}$$

Auxiliary distribution $q(\mathbf{z})$

Jensen's inequality

Convex functions



$$tf(x_1) + (1-t)f(x_2) \geq f(tx_1 + (1-t)x_2)$$

Generalized version
$$\int_x f(x)p(x) \geq f\left(\int_x xp(x)\right)$$

Maximizing a lower bound

$$\begin{aligned}\log p(\mathbf{x}) &= \log \int p(\mathbf{z}) p(\mathbf{x}|\mathbf{z}) d\mathbf{z} \\ &= \log \int q(\mathbf{z}) \frac{p(\mathbf{z})}{q(\mathbf{z})} p(\mathbf{x}|\mathbf{z}) d\mathbf{z} \\ &\geq \int q(\mathbf{z}) \log \left[\frac{p(\mathbf{z})}{q(\mathbf{z})} p(\mathbf{x}|\mathbf{z}) \right] d\mathbf{z}\end{aligned}$$

Auxiliary distribution $q(\mathbf{z})$

Jensen's inequality

Maximizing a lower bound

$$\log p(\mathbf{x}) = \log \int p(\mathbf{z}) p(\mathbf{x}|\mathbf{z}) d\mathbf{z}$$

$$= \log \int q(\mathbf{z}) \frac{p(\mathbf{z})}{q(\mathbf{z})} p(\mathbf{x}|\mathbf{z}) d\mathbf{z}$$

Auxiliary distribution $q(\mathbf{z})$

$$\geq \int q(\mathbf{z}) \log \left[\frac{p(\mathbf{z})}{q(\mathbf{z})} p(\mathbf{x}|\mathbf{z}) \right] d\mathbf{z}$$

Jensen's inequality

$$= \mathbb{E}_q \left[\log \frac{p(\mathbf{z})}{q(\mathbf{z})} \right] + \mathbb{E}_q [\log p(\mathbf{x}|\mathbf{z})]$$

Maximizing a lower bound

$$\log p(\mathbf{x}) \geq \mathbb{E}_q \left[\log \frac{p(\mathbf{z})}{q(\mathbf{z})} \right] + \mathbb{E}_q [\log p(\mathbf{x}|\mathbf{z})]$$

- **Reconstruction term**
 - Encourages the model to reconstruct the input
- If we parameterize $p(\mathbf{x}|\mathbf{z})$ as Gaussian

$$\begin{aligned} \log p(\mathbf{x}|\mathbf{z}) &= \log \mathcal{N}(\mathbf{x}; G_{\theta}(\mathbf{z}), \eta \mathbf{I}) \\ &= \log \left[\frac{1}{(2\pi\eta)^{D/2}} \exp \left(-\frac{1}{2\eta} \|\mathbf{x} - G_{\theta}(\mathbf{z})\|^2 \right) \right] \\ &= -\frac{1}{2\eta} \|\mathbf{x} - G_{\theta}(\mathbf{z})\|^2 + \text{const} \end{aligned}$$

Maximizing a lower bound

$$\log p(\mathbf{x}) \geq \mathbb{E}_q \left[\log \frac{p(\mathbf{z})}{q(\mathbf{z})} \right] + \mathbb{E}_q [\log p(\mathbf{x}|\mathbf{z})]$$

- Can be written as $-D_{\text{KL}}(q(\mathbf{z})\|p(\mathbf{z}))$. **KL term**.
- DKL is the Kullback-Leibler (KL) divergence $D_{\text{KL}}(q(\mathbf{z})\|p(\mathbf{z})) \triangleq \mathbb{E}_q \left[\log \frac{q(\mathbf{z})}{p(\mathbf{z})} \right]$
 - Widely used to measure the distance between two probability distributions
- Typically, $p(\mathbf{z}) = N(\mathbf{0}, \mathbf{1})$
 - The KL term encourages $q(\mathbf{z})$ to be close to the standard normal distribution $N(\mathbf{0}, \mathbf{1})$.

Maximizing a lower bound

Variational lower bound

$$\log p(\mathbf{x}) \geq \mathbb{E}_q [\log p(\mathbf{x}|\mathbf{z})] - D_{\text{KL}}(q\|p)$$

The role of each of the two terms:

The reconstruction term

$$\mathbb{E}_q[\log p(\mathbf{x}|\mathbf{z})] = -\frac{1}{2\sigma^2}\mathbb{E}_q[\|\mathbf{x} - G_{\theta}(\mathbf{z})\|^2] + \text{const}$$

is minimized when q is a **point mass** on

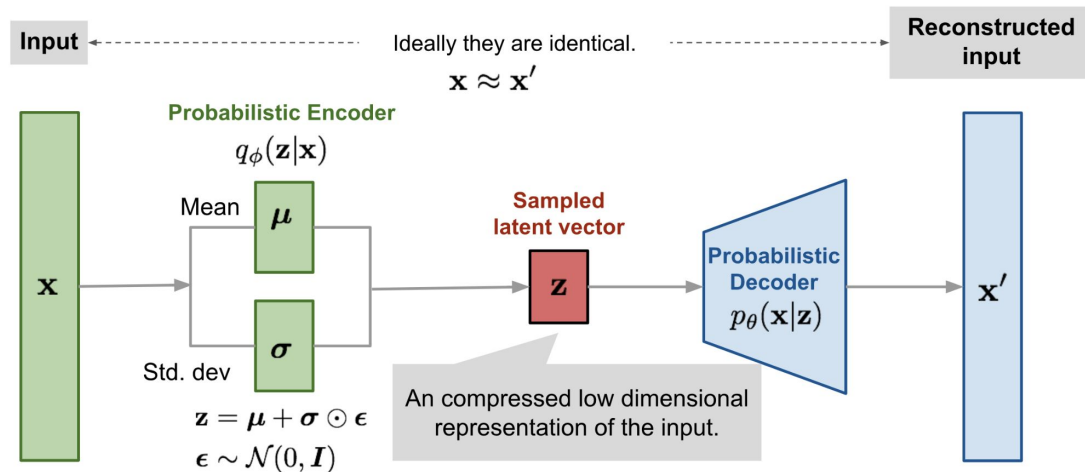
$$\mathbf{z}_* = \arg \min_{\mathbf{z}} \|\mathbf{x} - G_{\theta}(\mathbf{z})\|^2.$$

But a point mass would have infinite KL divergence.

VAE

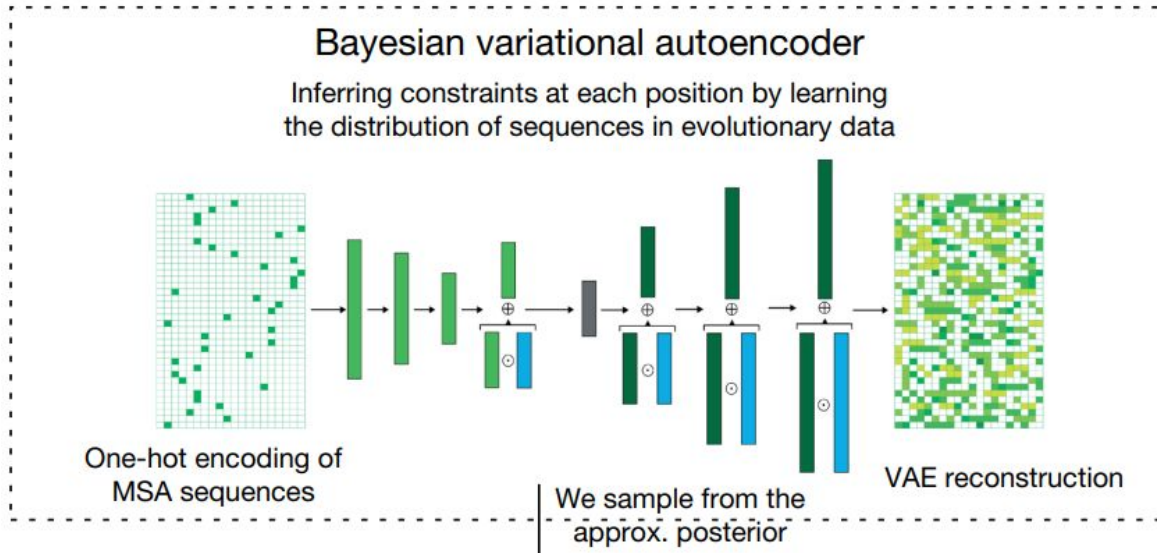
$$\log p(\mathbf{x}) \geq \mathbb{E}_q [\log p(\mathbf{x}|\mathbf{z})] - D_{\text{KL}}(q||p)$$

Further reading: “reparameterization trick”
(Kigima & Welling, 2013)



([image source](#))

Application: mutation effect prediction



↓


Evolutionary index

$$E_v \sim -\log \frac{P(x_v|\theta)}{P(x_{WT}|\theta)}$$


Approximating the negative log-likelihood ratio of mutant versus wild type

Application: molecule design

(a)



(b)



Summary of today

- PCA
 - Linear dimensionality reduction
 - Maximize variance
- Autoencoders
 - Nonlinear dimensionality reduction
 - Minimize reconstruction error
- VAE
 - Probabilistic generative model
 - Regularized latent space