

CSE8803/CX4803

Machine Learning in Computational Biology

Lecture 15:
Representation Learning in Graphs
(Network Embeddings)

Yunan Luo

Today's plan



- Graph representation learning
- Student presentation

Logistics



Questions posted on Ed during the weekend

Weisfeiler-Lehman kernel



- Assign initial colors



- Aggregate neighboring colors




- Hash aggregated colors




Hash table

1,1	-->	2
1,11	-->	3
1,111	-->	4
1,1111	-->	5



Colors
 $\phi(\text{Graph}) = [6, 2, 1, 2, 1, 0, 2, 1, 0, 0, 0, 2, 1]$
 Counts



Colors
 $\phi(\text{Graph}) = [6, 2, 1, 2, 1, 1, 1, 0, 1, 1, 1, 0, 1]$

Weisfeiler-Lehman kernel

- If the WL color assignments are **different**, the two graphs are **not** isomorphic
- If the WL color assignments are the **same**, are the two graphs isomorphic?
 - possibly, but not necessarily

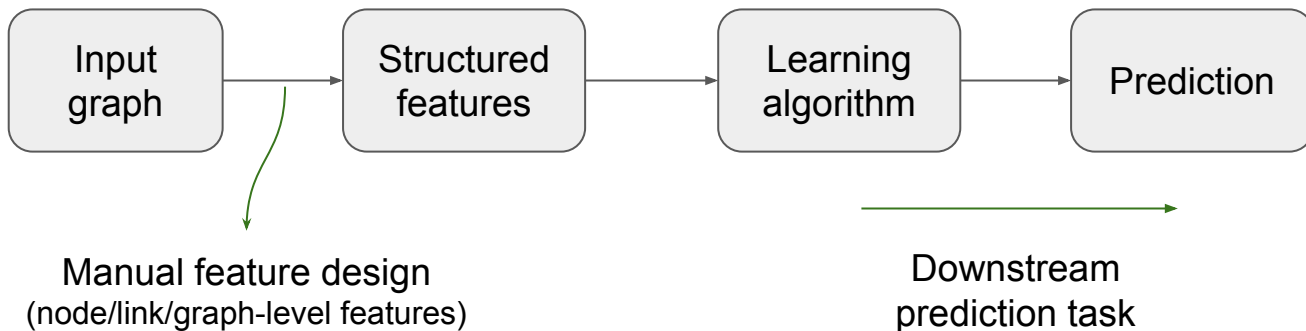


https://blog.twitter.com/engineering/en_us/topics/insights/2021/provably-expressive-graph-neural-networks

- How many iterations needed for the color refinement?
 - To test the isomorphism between two graphs, N (#nodes) iterations
 - Feature vectors, can choose a predefined step $h \leq N$

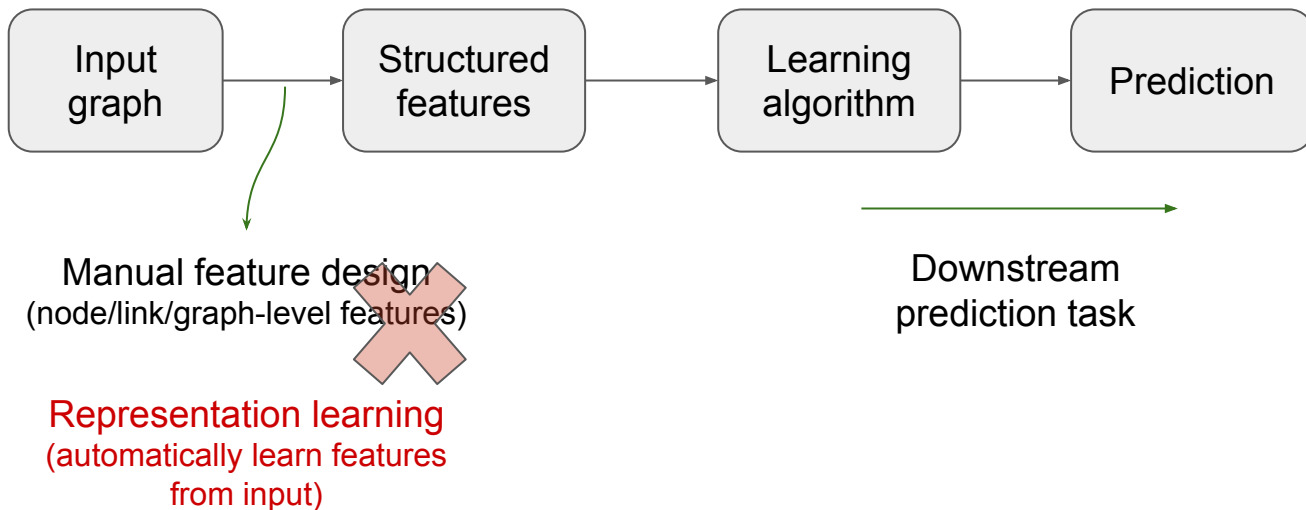
Last lecture: traditional ML for graphs

- Given an input graph, extract node, link and graph-level features, learn a model (SVM, neural network, etc.) that maps features to labels.




Graph representation learning

- Graph representation learning: learn a feature for each node from the graph input automatically



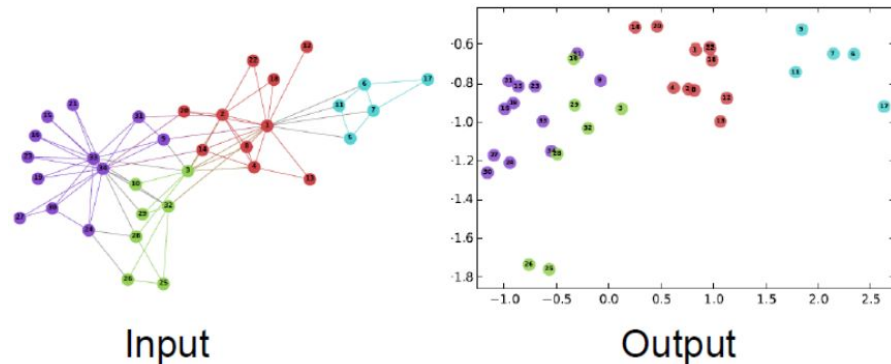
Graph representation learning

- **Goal:** Efficient task-independent feature learning for machine learning with graphs




Why embedding?

- Similarity of embeddings between nodes indicates their similarity in the network.



DeepWalk: Online Learning of Social Representations. KDD 2014

- Potentially used for many downstream predictions




Tasks

- Node classification
- Link prediction
- Graph classification
- Anomalous node detection
- Clustering
-

Embedding nodes


- **Goal:** encode nodes so that **similarity in the embedding space** (e.g., dot product) approximates **similarity in the graph**




Embedding nodes

$$\underset{\text{In the original network}}{\text{similarity}(u, v)} \approx \underset{\text{Similarity of the embeddings}}{\mathbf{z}_v^T \mathbf{z}_u}$$

$\text{ENC}(v)$: an **encoder** maps from
nodes to embeddings



Embedding nodes



Two key components

- **Encoder**: maps each node to a low-dimensional vector

$$\text{ENC}(\mathbf{v}) = \mathbf{Z}_v$$

Node in the input graph

d -dimensional vector (embedding)

- **Similarity function**: specifies how the relationships in vector space map to the relationships in the original network

$$\text{similarity}(u, v) \approx \mathbf{Z}_v^T \mathbf{Z}_u$$

Similarity of u and v in the original network


Dot product between node embeddings

How to define node similarity in the network?

- Key choice of methods is how to **define node similarity**.
- Should two nodes have a similar embedding if they ...
 - are linked?
 - share neighbors?
 - ...
- This lecture: define node similarity based “**topological roles**” of each node with respect to other nodes.
- Two graph representation learning algorithms:
 - Diffusion component analysis [DCA] (Cho et al, 2016, *Cell Systems*)
 - Node2vec (Grover et al, 2016, *KDD*)


Motivating example

Example: protein function prediction




Voting by direct neighbors

If there is no direct neighbor with known function




Shortest path




Floyd-Warshall algorithm: all pairwise distances
Computational Complexity: $O(n^3)$

Is shortest path a good metric?



Label propagation algorithm



**How to solve
this problem?**

Connected nodes tend to have similar function (color).

$$\min_{f_{red}} \sum_{(i,j) \in E} w_{ij} (f_{red}(i) - f_{red}(j))^2$$


$$\forall i \in RED, f_{red}(i) = 1$$

Diffusion component analysis (DCA)

[1] Cho, Hyunghoon, Bonnie Berger, and Jian Peng. "Diffusion component analysis: unraveling functional topology in biological networks." *International Conference on Research in Computational Molecular Biology*. Springer, Cham, 2015.

[2] Cho, Hyunghoon, Bonnie Berger, and Jian Peng. "Compact integration of multi-network topology for functional analysis of genes." *Cell systems* 3.6 (2016): 540-548.

Random walk and Pagerank



Article

Talk

Re

PageRank

From Wikipedia, the free encyclopedia


PageRank (PR) is an [algorithm](#) used by [Google Search](#) to rank [web pages](#) in their search [engine results](#). It is named after both the term "web page" and co-founder [Larry Page](#).

PageRank is a way of measuring the importance of website pages. According to Google:

PageRank works by counting the number and quality of links to a page to determine a rough estimate of how important the website is. The underlying assumption is that more important websites are likely to receive more links from other websites.^[1]

Random walk with restart

Start from node i




Initialization

$$s_i = (0, 0, \dots, 1, \dots, 0)$$

Random walk with restart

Distribute to neighbors




four neighbors

Propagation

$$s_i = (0, \dots, 0.25, \dots, 0.25, \dots, 0.25, \dots, 0.25, \dots, 0)$$

Random walk with restart


Shrinkage and restart from node i




Restart

$$s_i = (0, \dots, 0.125, \dots, 0.125, \dots, 0.5, \dots, 0.125, \dots, 0.125, \dots, 0)$$

Random walk with restart



Random walk with restart



Random walk with restart

Adjacency matrix $A : A_{ij}$

Transition matrix $B \in R^{n \times n} : B_{ij} = A_{ij} / \sum_k A_{ik}$

Restart probability p

Algorithm: Repeat

$$\forall i, j \quad s_i(j) = (1 - p) \sum_k s_i(k) B_{kj} + p \delta(i = j)$$

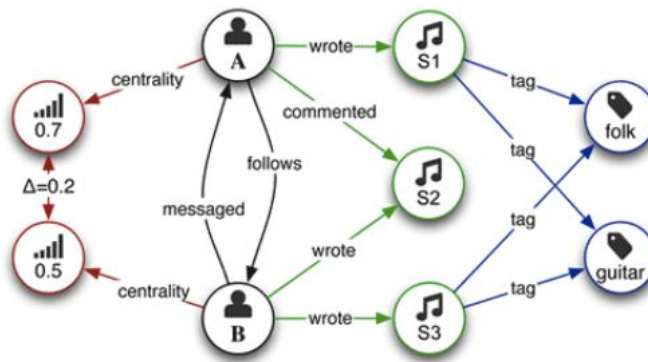

Matrix operation

$$S^{new} = (1 - p) S^{old} B + p I_n$$

Random walk with restart

- **Comments:**

- Very simple implementation
- Capture long-range relationship in the graph
- Robust to missing edges
- Many applications in social network analysis, web data analysis, and bioinformatics



Weighted Voting

Assign color vector:

$$\forall i \in RED, f_{red}(i) = 1$$
$$\forall i \notin RED, f_{red}(i) = 0$$

Majority voting

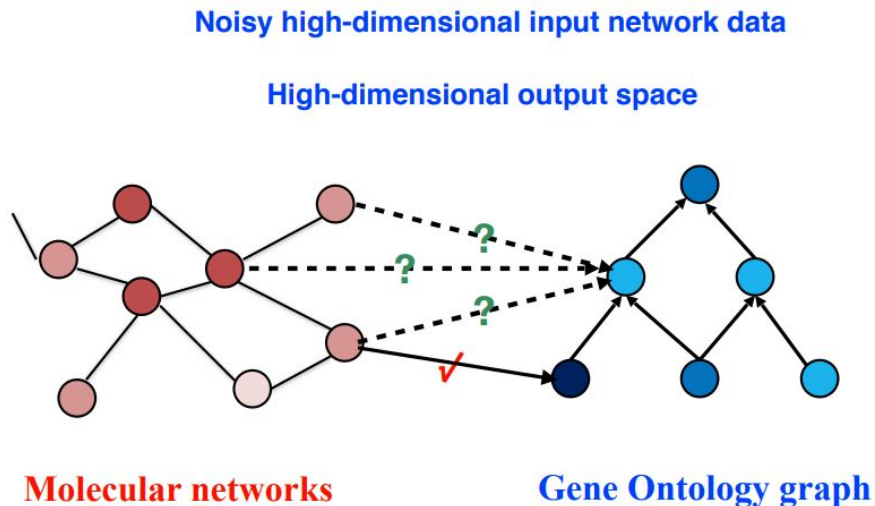
$$p_{red}(i) = \sum_j Sim(i, j) f_{red}(j)$$

Some good similarity metrics:


1. reciprocal of shorted distance $1/d(i, j)$
2. random walk similarity $s_i(j) + s_j(i)$
3. diffusion state similarity $1/\|s_i - s_j\|_1$

Supervised learning: protein function prediction

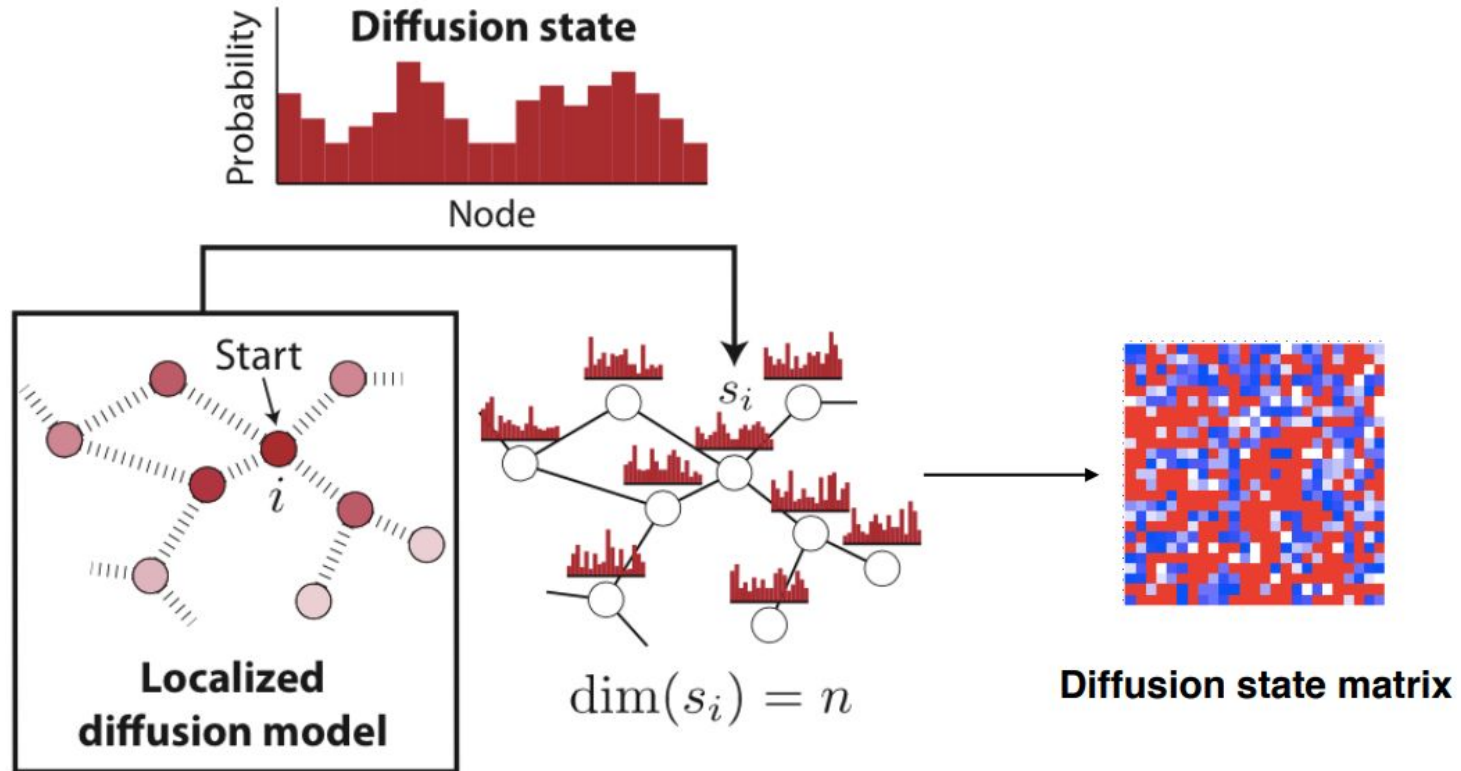
- More than 16K proteins in human protein-protein interaction network
- There may be many **fake** and **missing** edges in the network
- Nearly **half of** the functions only have **very few** annotations
- **Not enough** training samples for these functions
- A simple classifier will be **overfitting**



Dimensionality reduction




After we run random walk with restart



We can run SVD or NMF on S

- Let us assume that we don't know how this diffusion state matrix was generated



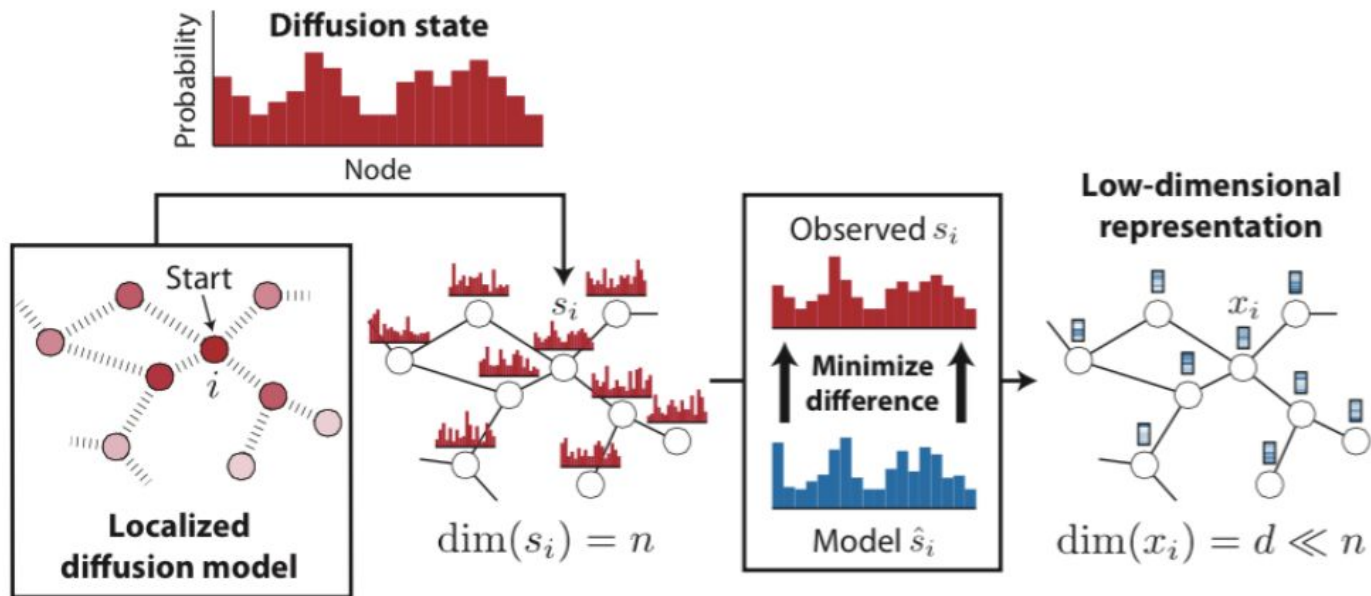
Any issues with these dimensionality reduction algorithms?

$$\min_{X, W} \|S - XW\|_2^2$$

$$X \in R^{n \times d}, W \in R^{d \times n}$$


Diffusion component analysis (DCA)

- **Idea:** we hope to construct a model to **approximate** observed diffusion states



After random walk (diffusion)


- $s_i(j)$ gives the probability of reaching node j after the random walk from node i



$$\sum_j s_i(j) = 1$$
$$\forall i, j \quad s_i(j) \geq 0$$

What would be a good way to model diffusion states?


Logistic regression



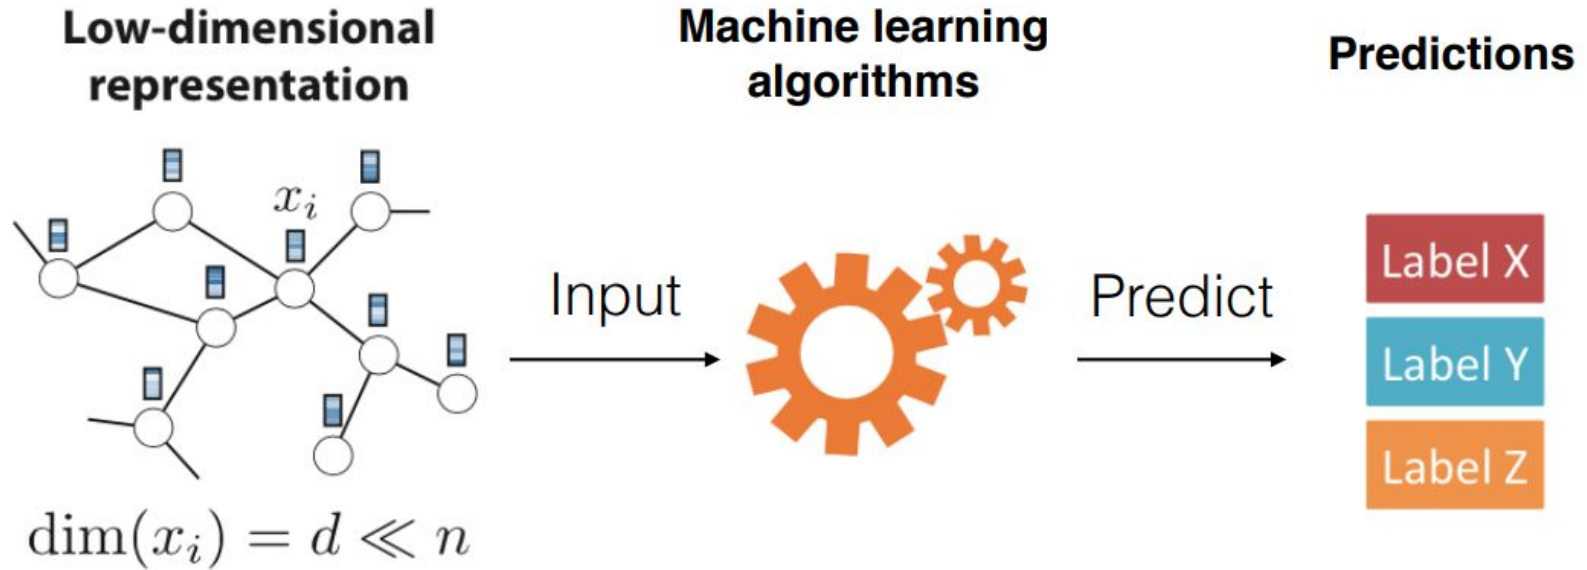
$$\hat{s}_i(j) = \frac{\exp(w_i^T x_j)}{\sum_k \exp(w_i^T x_k)}$$

Is it better than SVD/NMF?


Matching observed and model data



Machine learning with network data



Comparisons



Analogy to PCA



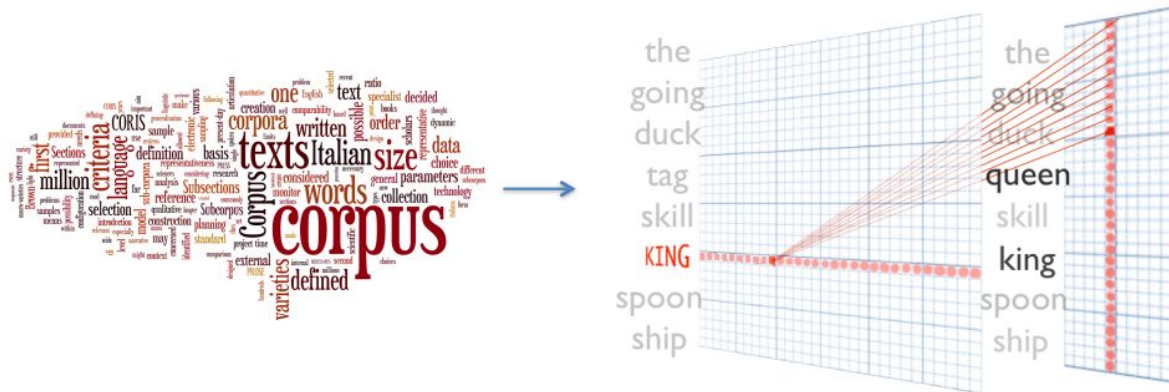
Principal component analysis (PCA)

- Input: **matrix** data
- Goal: find low-rank approximation that best explain the **variance** of the matrix input

Diffusion component analysis (DCA)

- Input: **network** data
- Goal: find low-dim representations that best explain the **topological pattern** of the network input


Similarity to word2vec



If we have many networks

Step 1: run diffusion within each network

Step 2: jointly optimize vectors over all networks




Shared vectors capture global patterns

node2vec

Grover et al. “node2vec: Scalable Feature Learning for Networks”. *KDD*, 2016


node2vec

- **Goal:** Embed nodes with similar network neighborhoods close in the feature space
- **Idea:** Use flexible, **biased random walks** that can trade off between **local** and **global** views of the network



Biased walks

Two classic strategies to define a neighborhood $N(u)$ of a given node u



Walk of length 3 ($N(u)$ of size 3):


$$N_{BFS}(u) = \{s_1, s_2, s_3\}$$

Local microscopic view

$$N_{DFS}(u) = \{s_4, s_5, s_6\}$$


Global microscopic view

BFS & DFS



BFS:

Micro-view of
neighbourhood



DFS:

Macro-view of
neighbourhood

How to interpolate BFS & DFS?

Biased fixed-length random walk \mathbf{R} that, given a node \mathbf{u} , generates neighborhood $\mathbf{N}_R(\mathbf{u})$


Two parameters:

- Return parameter \mathbf{p} :
 - Return back to the previous node
- In-out parameter \mathbf{q} :
 - Moving outwards (DFS) vs. inwards (BFS)
 - Intuitively, \mathbf{q} is the “ratio” of BFS vs. DFS

Biased random walk


Biased 2nd-order random walks explore network neighborhoods

- Random walk just traversed edge (s_1, w) and is now at w
- **Observation:** Neighbors of w can only be



Biased random walk

Walker came over edge (s_1 , w) and is now at w . Where to go next?




$1/p$, $1/q$, and 1 are
unnormalized probabilities

p , q model transition probabilities

- p : return parameter
- q : “walk away” parameter

Biased random walk

Walker came over edge (s_1, w) and is now at w . Where to go next?



$w \rightarrow$

Target t	Prob.	Dist. (s_1, t)
s_1	$1/p$	0
s_2	1	1
s_3	$1/q$	2
s_4	$1/q$	2

$1/p$, $1/q$, and 1 are
unnormalized probabilities

- **DFS-like walk:** High value of $1/p$
- **BFS-like walk:** High value of $1/q$
- **Random walk $N_R(u)$:** nodes visited by the biased walk

Representation learning framework

- Given $G = (V, E)$
- Goal:** learn a mapping $f: u \rightarrow R^d: f(u)=z_u$
- Intuition:** learn representations such that given the representation z_u of node u , we can predict what its neighbors $N_R(u)$ are
- Log-likelihood objective**

$$\max_f \sum_{u \in V} \log P(N_R(u) | z_u) \quad \leftarrow \text{Maximum likelihood objective}$$

$N_R(u)$ is the neighborhood of node u by strategy R

Optimization

- Log-likelihood objective

$$\max_f \sum_{u \in V} \log P(N_R(u) | \mathbf{z}_u)$$

- Equivalently

$$\mathcal{L} = \sum_{u \in V} \sum_{v \in N_R(u)} -\log(P(v | \mathbf{z}_u))$$

- **Intuition**: maximize the similarity between \mathbf{u} and other nodes in the walk $N_R(\mathbf{u})$
- Parameterize $P(\mathbf{v} | \mathbf{z}_u)$ using softmax

$$P(v | \mathbf{z}_u) = \frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)}$$

Optimization

- Putting it all together

$$\mathcal{L} = \sum_{u \in V} \sum_{v \in N_R(u)} -\log\left(\frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)}\right)$$

- Finding representations \mathbf{z}_u that minimize \mathcal{L}

Optimization

$$\mathcal{L} = \sum_{u \in V} \sum_{v \in N_R(u)} -\log\left(\frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)}\right)$$

Expensive to compute.
Need approximation.

- **Solution:** Negative sampling

$$\log\left(\frac{\exp(\mathbf{z}_u^T \mathbf{z}_v)}{\sum_{n \in V} \exp(\mathbf{z}_u^T \mathbf{z}_n)}\right) \approx \log(\sigma(\mathbf{z}_u^T \mathbf{z}_v)) - \sum_{i=1}^k \log(\sigma(\mathbf{z}_u^T \mathbf{z}_{n_i})), n_i \sim P_V$$

Sigmoid function

(Makes each term a
"probability" in $[0, 1]$)

Random distribution
over nodes

- **Idea:** Instead of normalizing w.r.t all nodes, just normalize against k random negative samples
- Sample k negative nodes each with prob. proportional to its degree

node2vec algorithm

1. Compute random walk probabilities
 2. Simulate r random walks of length l starting from each node u
 3. Optimize the node2vec objective using Stochastic Gradient Descent
- Linear-time complexity
 - All 3 steps are individually parallelizable

Other random walk-based methods

- Different kinds of biased random walks:
 - Based on node attributes ([Dong et al., 2017](#))
 - Based on learned weights ([Abu-El-Haija et al., 2017](#))
- Alternative optimization schemes:
 - Directly optimize based on 1-hop and 2-hop random walk probabilities (e.g., LINE from [Tang et al. 2015](#))
- Network preprocessing techniques:
 - Run random walks on modified versions of the original network (e.g., [Ribeiro et al. 2017's struct2vec](#), [Chen et al. 2016's HARP](#))

Summary of today

- Representations learning in graphs
 - Learning embeddings that capture structure/topological similarity between network nodes
- Unsupervised representation learning framework
 - Random-walk based
- Two algorithms
 - Diffusion Component Analysis (DCA)
 - Node2vec